

SAI ABHIRAM GOUD SEEKATI

Abhiramsgoud2003@gmail.com | +1 (361) 695-4112 | [LinkedIn](#) | United States (Open to Relocation) | [GitHub](#) | [Portfolio](#)

SUMMARY

AWS Certified Data Engineer (Associate) and Cloud Practitioner with hands-on experience building end-to-end data pipelines, machine learning systems, and BI platforms across cloud-native and on-premise environments. Delivered a 35% improvement in ETL throughput across 500K+ records and reduced ad-hoc reporting requests by 40% through Power BI dashboards powered by dimensional data models and advanced DAX — patterns directly aligned with AWS Glue, Amazon RDS, and Amazon QuickSight. Engineered a full TensorFlow/Keras Autoencoder fraud detection pipeline achieving 15% accuracy improvement and 12% reduction in false positives, with architecture consistent with Amazon SageMaker end-to-end ML workflows. Proficient in Python, SQL, PostgreSQL, Power BI, TensorFlow, Scikit-learn, Flask, Docker, and AWS services including S3, Glue, SageMaker, Athena, Lambda, and MWAA. Pursuing M.S. in Computer Science at Texas A&M University – Corpus Christi; open to full-time Data Analyst, Data Engineer, or ML Engineer roles.

EDUCATION

Texas A&M University – Corpus Christi, TX

Aug 2024 – May 2026

Master of Science, Computer Science

Guru Nanak Institute of Technology – Hyderabad, India

Aug 2020 – May 2024

Bachelor of Engineering, Computer Science (AI & ML)

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript, R, Bash/Shell Scripting.

ML & AI: TensorFlow, Keras, Scikit-learn, LSTM neural networks, NLP, Logistic Regression, Anomaly Detection, Generative AI, Feature Engineering, A/B Testing.

Data Engineering: ETL/ELT Pipelines, PostgreSQL, MySQL, SQLite, DuckDB, MongoDB, Data Modeling, Data Governance, Metadata Management.

Systems & Infra: Linux/Bash, Distributed & Parallel Computing (MPI), Docker, AWS S3, Git/GitHub, REST APIs, Flask, React.

BI & Visualization: Power BI (DAX, Power Query), Excel (Pivot Tables, Advanced Formulas), Matplotlib, Seaborn.

Methods: Agile/Scrum, Statistical Analysis, EDA, Data Validation, Software Development Lifecycle, Documentation.

PROFESSIONAL EXPERIENCE

Data Analyst Intern | SAS Data and AI Solutions

Aug 2025 – Dec 2025

- Developed and optimized Python-based ETL pipelines against PostgreSQL — following patterns consistent with AWS Glue and Amazon RDS — processing 500K+ energy and billing records with indexing, CTEs, and execution plan analysis to improve throughput by 35%.
- Engineered interactive dashboards using dimensional data models and advanced calculated measures in Power BI, reducing ad-hoc reporting requests by 40% and accelerating executive decision cycles; directly applicable to Amazon QuickSight with SPICE and custom calculated fields.
- Implemented data governance workflows including metadata management, validation, reconciliation, and root cause investigation scripts in Python — mirroring AWS Glue Data Catalog and Amazon DataZone governance patterns — reinforcing KPI reliability across enterprise systems.
- Collaborated cross-functionally to gather business and technical requirements, producing data mapping specifications and interface documentation that enabled seamless system integrations, equivalent to designing data contracts for AWS Glue and Amazon AppFlow pipelines.
- Contributed to reporting pipeline automation using Python scripting in a Linux/Unix environment, improving reliability and reducing manual effort — skills directly transferable to AWS Lambda functions and Amazon MWAA DAG authoring.

Data Analyst Intern | Campalin Innovations

Jan 2023 – Aug 2023

- Designed and trained an Autoencoder neural network in TensorFlow/Keras to detect anomalies across 500K+ credit card transactions, engineering the full model pipeline from data ingestion to inference — a workflow directly transferable to Amazon SageMaker for end-to-end model training, hosting, and inference endpoints — achieving a 15% improvement in fraud detection accuracy.
- Applied SMOTE-based data augmentation and feature engineering using Python, Pandas, and NumPy to address class imbalance in transaction data, reducing false positives by 12% and improving model reliability for production deployment; aligns with Amazon SageMaker Data Wrangler and SageMaker Experiments for preprocessing and iterative feature development.
- Conducted large-scale EDA using Pandas, Matplotlib, and Seaborn to uncover 3 critical fraud behavior patterns, translating findings into algorithmic improvements that directly enhanced detection performance — skills applicable to data profiling on Amazon Athena and visual analysis through Amazon QuickSight.
- Optimized the model training pipeline through vectorized operations and batch processing, cutting end-to-end execution time by 20% and enabling faster experimentation cycles — consistent with SageMaker Pipelines for automated workflow orchestration and SageMaker Training Jobs with distributed compute configurations.

PROJECTS

Intelligent Customer Support System | NLP, Flask, PostgreSQL, Python, Explainable AI | [GitHub](#)

- Architected an end-to-end NLP-powered chatbot system with a RESTful Flask backend and PostgreSQL database (handling structured customer order data), reducing average query resolution time through automated intent classification and entity extraction.
- Implemented explainable AI (XAI) techniques to surface model reasoning behind chatbot responses, improving transparency and decision interpretability by providing confidence scores alongside predictions.
- Designed and deployed the full software stack from schema design and ETL pipelines to a React frontend (port 5173) communicating with the Flask API — demonstrating end-to-end software development lifecycle ownership.
- Engineered a relational data model in PostgreSQL supporting multi-turn conversation history, order lookups, and session state management, enabling scalable distributed query execution across concurrent user sessions.

Revenue Leakage & Customer Churn Intelligence Platform | [GitHub](#)

- Built an end-to-end churn and revenue-leakage analytics platform integrating subscription, invoice, usage, and support data using SQL queries against relational databases and Power BI.
- Cleaned, validated, and modeled data from multiple sources; developed an explainable churn-scoring model (logistic regression, Scikit-learn) and surfaced risk tiers in Power BI dashboards with drill-through views.

Demand Shock Forecast Resilience Analytics Platform | [GitHub](#)

- Built a forecasting analytics platform using Python, SQLite, and Power BI, generating 30-day forecasts and 72,000 scenario impact records across multi-store retail demand.
- Implemented scenario-based stress testing and created executive dashboards for forecast quality monitoring and SKU/store prioritization, translating data into business-ready insights.

CERTIFICATIONS

AWS Certified Data Engineer - Associate certificate · AWS Certified Cloud Practitioner (CLF-C02) · Alteryx Masterclass for Data Analytics · ETL and Reporting Certificate · Data Science Internship Certificate · Google Analytics Certificate · SQL Certificate. Full list: [here](#). Badges: [here](#)

LANGUAGES

English (Fluent) · Telugu (Native) · Hindi (Fluent)